

Adarsh Solanki

As5nr

CS 2150: In-Lab 9

4/11/12

Optimized Code

For this section of the lab, I compiled a simple C file (tried C++, but the generated assembly was identical for both according to diff because no C++ specific features were used)

```
1 int main() {  
2  
3     int x = 5 + 4; // 9  
4     int y = 3 - 2; // 1  
5  
6     int z = (x + y)/2; // 5  
7     z = 4;  
8     return z;  
9 }
```



```
1 .section __TEXT,__text,regular  
2 .globl _main  
3 .align 4, 0x90  
4 _main:  
5 Lehfnc_begin1:  
6 pushq %rbp  
7 Ltmp0:  
8 movq %rsp, %rbp  
9 Ltmp1:  
10 movl $9, -12(%rbp)  
11 movl $1, -16(%rbp)  
12 movl -12(%rbp), %eax  
13 movl -16(%rbp), %ecx  
14 addl %ecx, %eax  
15 movl %eax, %ecx  
16 shrl $31, %ecx  
17 leal (%rax,%rcx), %eax  
18 sarl %eax  
19 movl %eax, -20(%rbp)  
20 movl $4, -20(%rbp)  
21 movl $4, -8(%rbp)  
22 movl $4, -4(%rbp)  
23 movl -4(%rbp), %eax  
24 popq %rbp  
25 ret  
26 Lehfnc_end1:
```

Above: original C code

Right: basic g++ assembly output

*Below: optimized (O2) g++
assembly output*



```
1 .section __TEXT,__text,regular  
2 .globl _main  
3 .align 4, 0x90  
4 _main:  
5 Lehfnc_begin1:  
6 pushq %rbp  
7 Ltmp0:  
8 movq %rsp, %rbp  
9 Ltmp1:  
10 movl $4, %eax  
11 popq %rbp  
12 ret  
13 Lehfnc_end1:
```

The difference with the -O2 flag is tremendous! The entire main function was a 17 instruction x86 segment that closely matched the source C code. There are definitely also some inefficiencies, for example at the very end, there are a lot of extraneous 4's being placed onto the stack and then eventually one is pushed into the return register and the function returns.

Conversely, on the optimized assembly, g++ did whatever it could to minimize the number of instructions required to run the output assembly, so it took the time to pre-process things like the return value. The main function is 5 instructions now, 4 of which are standard boilerplate like saving/restoring the base pointer and calling ret to return. The way I see it, the O2 flag gives g++ the directive to take extra time now to make it more efficient later, pre-compute anything that you already have enough information to compute and cut-out any unnecessary movement/storage of data.

Cout Function

The next type of optimization I explored was optimization of I/O using standard functions: namely cout. Although, I didn't expect to really be able to understand most of the generated assembly, as I imagine I/O to be complex on the lower levels, I did gain some valuable insights throughout reading through the assembled programs. I ended up reading a good deal into the differences between the standard x86-64 architecture and that of x86-64 Mac OS X computers.

The following are screen captures of the C++ source, the default assembly, and the optimized assembly (in that order):

```
Vim
1 #include <iostream>
2
3 using namespace std;
4
5 int main() {
6     cout << 127 << endl;
7 }
```

Original Source

Truncated assembly

```
Vim
1 .section __TEXT,__text,regular,pure_instructions
2 .globl _main
3 .align 4,0x90
4 _main:
5 Leh_func_begin1:
6     pushq %rbp
7 Ltmp0:
8     movq %rsp,%rbp
9 Ltmp1:
10    subq $16,%rsp
11 Ltmp2:
12    movq __ZSt4cout@GOTPCREL(%rip),%rax
13    leaq (%rax),%rax
14    movl $127,%ecx
15    movq %rax,%rdi
16    movl %ecx,%esi
17    callq __ZNSolsEi
18    movq __ZSt4endlIcSt11char_traitsIcEERSt13basic_ostreamIT_0_ES6_@GOTPCREL(%rip),%rcx
19    leaq (%rcx),%rcx
20    movq %rax,%rdi
21    movq %rcx,%rsi
22    callq __ZNSolsEPFRSoS_E
23    movl $0,-8(%rbp)
24    movl -8(%rbp),%eax
25    movl %eax,-4(%rbp)
26    movl -4(%rbp),%eax
27    addq $16,%rsp
28    popq %rbp
29    ret
30 Leh_func_end1:
31
32 .section __TEXT,__StaticInit,regular,pure_instructions
33 .align 4,0x90
34 __GLOBAL__I_main:
35 Leh_func_begin2:
36     pushq %rbp
37 Ltmp3:
38     movq %rsp,%rbp
39 Ltmp4:
40     movl $1,%eax
41     movl $65535,%ecx
42     movl %eax,%edi
43     movl %ecx,%esi
44     callq __Z41__static_initialization_and_destruction_0ii
45     popq %rbp
46     ret
47 Leh_func_end2:
48
49 .align 4,0x90
50 __Z41__static_initialization_and_destruction_0ii:
51 Leh_func_begin3:
52     pushq %rbp
53 Ltmp5:
54     movq %rsp,%rbp
55 Ltmp6:
```

Optimized Assembly

```
1  .section __TEXT,__text,regular,pure_instructions
2  .globl _main
3  .align 4, 0x90
4  _main:
5  Leh_func_begin1:
6  pushq %rbp
7  Ltmp0:
8  movq %rsp, %rbp
9  Ltmp1:
10 movq __ZSt4cout@GOTPCREL(%rip), %rdi
11 movl $127, %esi
12 callq __ZNSolsEi
13 movq __ZSt4endlC@GOTPCREL(%rip), %rsi
14 movq %rax, %rdi
15 callq __ZNSolsEPFRSoS_E
16 xorl %eax, %eax
17 popq %rbp
18 ret
19 Leh_func_end1:
20
21 .section __TEXT,__StaticInit,regular,pure_instructions
22 .align 4, 0x90
23 __GLOBAL__I_main:
24 Leh_func_begin2:
25 pushq %rbp
26 Ltmp2:
27 movq %rsp, %rbp
28 Ltmp3:
29 leaq __ZStL8__ioinit(%rip), %rdi
30 callq __ZNSt8ios_base4InitC1Ev
31 leaq __tcf_0(%rip), %rdi
32 xorl %esi, %esi
33 movq __dso_handle@GOTPCREL(%rip), %rdx
34 popq %rbp
35 jmp __cxa_atexit # TAILCALL
36 Leh_func_end2:
37
38 .section __TEXT,__text,regular,pure_instructions
39 .align 4, 0x90
40 __tcf_0:
41 Leh_func_begin3:
42 pushq %rbp
43 Ltmp4:
44 movq %rsp, %rbp
45 Ltmp5:
46 leaq __ZStL8__ioinit(%rip), %rdi
47 popq %rbp
48 jmp __ZNSt8ios_base4InitD1Ev # TAILCALL
49 Leh_func_end3:
50
51 .zerofill __DATA,__bss,__ZStL8__ioinit,1,3
52 .section __DATA,__mod_init_func,mod_init_funcs
53 .align 3
54 .quad __GLOBAL__I_main
55 .section __TEXT,__eh_frame,coalesced,no_toc+strip_static_syms+live_support
56 EH_frame0:
57 Lsection_eh_frame:
```

The first thing I noticed is that the optimized assembly code passes much less to the `__ZNSolsEi` function that is called. The original code passes in `__ZSt4cout@GOTPCREL($rip)` [via `$RAX`] to the function. This intrigued me, as it seemed like it was passing in a function to the function.

It turns out, the GOTPCREL(\$rip) is the Mac x86-64 environment's way of accessing local and small data. (Source:

http://developer.apple.com/library/mac/documentation/DeveloperTools/Conceptual/MachOTopics/1-Articles/x86_64_code.html)

My interpretation of this is that the prefix __ZSt4cout is a label for one of the functions that exists in the iostream library (the cout function, perhaps a specific one to print a number vs. printing another data type) and is represented in the program space's "global offset table" (GOT) which uses "RIP-relative addressing" which means addressing relative to the instruction pointer.

The optimized code performs very similar things, but in fewer instructions. For example, the default code stores the ZSt4cout function pointer in \$rax and then moves it from \$rax to \$rdi. This is probably due to convention. The optimized code on the other hand stores the information directly into \$rdi. I don't understand why modern compilers don't always optimize little steps like these, as it doesn't even seem to be an optimization, but rather an improvement on the inefficient assembly produced by default. I'm sure there is a reason for this that involves things beyond the scope of my understanding.

Control Flow

The final optimization that I tried to study was that of control flow. I made a C++ program that included a very primitive demonstration of if/else and a while loop.



```
1 #include <iostream>
2
3 using namespace std;
4
5 int main() {
6     int x = 10;
7     if (x > 5) {
8         while ( x > 5 ) {
9             x--;
10        }
11    } else {
12        x += 8;
13    }
14 }
```

```

1  [section __TEXT,__text,
2  .globl _main
3  .align 4,0x90
4  _main:
5  Leh_func_begin1:
6  pushq %rbp
7  Ltmp0:
8  movq %rsp,%rbp
9  Ltmp1:
10 movl $10,-12(%rbp)
11 movl -12(%rbp),%eax
12 cmpl $5,%eax
13 jle LBB1_5
14 jmp LBB1_3
15 LBB1_2:
16 movl -12(%rbp),%eax
17 subl $1,%eax
18 movl %eax,-12(%rbp)
19 LBB1_3:
20 movl -12(%rbp),%eax
21 cmpl $5,%eax
22 jg LBB1_2
23 jmp LBB1_6
24 LBB1_5:
25 movl -12(%rbp),%eax
26 addl $8,%eax
27 movl %eax,-12(%rbp)
28 LBB1_6:
29 movl $0,-8(%rbp)
30 movl -8(%rbp),%eax
31 movl %eax,-4(%rbp)
32 movl -4(%rbp),%eax
33 popq %rbp
34 ret
35 Leh_func_end1:
36

```

default

```

1  [section __TEXT,__text,regular,pure_instructions
2  .globl _main
3  .align 4,0x90
4  _main:
5  Leh_func_begin1:
6  pushq %rbp
7  Ltmp0:
8  movq %rsp,%rbp
9  Ltmp1:
10 xorl %eax,%eax
11 popq %rbp
12 ret
13 Leh_func_end1:
14
15 [section __TEXT,__StaticInit,regular,pure_instructions
16 .align 4,0x90
17 __GLOBAL__I_main:
18 Leh_func_begin2:
19 pushq %rbp
20 Ltmp2:
21 movq %rsp,%rbp
22 Ltmp3:
23 leaq __ZStL8_iocinit(%rip),%rdi
24 callq __ZNSt8ios_base4InitC1Ev
25 leaq __tcf_0(%rip),%rdi
26 xorl %esi,%esi
27 movq __dso_handle@GOTPCREL(%rip),%rdx
28 popq %rbp
29 jmp __cxa_atexit # TAILCALL
30 Leh_func_end2:
31
32 [section __TEXT,__text,regular,pure_instructions
33 .align 4,0x90
34 __tcf_0:
35 Leh_func_begin3:
36 pushq %rbp
37 Ltmp4:
38 movq %rsp,%rbp
39 Ltmp5:
40 leaq __ZStL8_iocinit(%rip),%rdi
41 popq %rbp
42 jmp __ZNSt8ios_base4InitD1Ev # TAILCALL
43 Leh_func_end3:

```

optimized

The loop iteration in the non-optimized code is pretty similar to how I would code it myself. It creates a few labels for each of the control flow “states” and then uses the basic `cmp`, `jle`, `j`, `jge` instructions that I would use, were I to code this myself.

The optimized code on the other hand, is very different. It is much harder to draw similarities between the C++ code and the optimized assembly code. I tried to search for some common instructions like `cmp` or any jump, but to no avail, so I decided to dig around in the code and see what I could find through the Internet.

The first thing that jumped out to me was the line

```
jmp __cxa_atexit # TAILCALL
```

This seemed odd, because I hadn’t yet seen g++ generate comments in assembled code. It turned out that this was actually present in the default assembly code, as well as every other assembled program I had generated with g++. It turned out to be g++’s way of doing constructors, so this must have been

some sort of initialization of the main() object? This leads me to believe that this program simply does nothing, as the first (main) function simply xor's EAX with itself (zero's it) and then returns. The compiler rendered my function equivalent to "return 0;" and come to think of it, it technically was, as none of the variables or register values have any importance outside of the scope of the function itself.