

Adarsh Solanki

As5nr

CS 2150: In-Lab 9

4/11/12

Optimized Code

For this section of the lab, I compiled a simple C file (tried C++, but the generated assembly was identical for both according to diff because no C++ specific features were used)

```
1 int main() {
2
3     int x = 5 + 4; // 9
4     int y = 3 - 2; // 1
5
6     int z = (x + y)/2; // 5
7     z = 4;
8     return z;
9 }
```



```
1 .section __TEXT,__text,regular
2 .globl _main
3 .align 4, 0x90
4 _main:
5 Leh_func_begin1:
6     pushq    %rbp
7 Ltmp0:
8     movq     %rsp, %rbp
9 Ltmp1:
10    movl     $9, -12(%rbp)
11    movl     $1, -16(%rbp)
12    movl     -12(%rbp), %eax
13    movl     -16(%rbp), %ecx
14    addl     %ecx, %eax
15    movl     %eax, %ecx
16    shrl     $31, %ecx
17    leal     (%rax,%rcx), %eax
18    sarl     %eax
19    movl     %eax, -20(%rbp)
20    movl     $4, -20(%rbp)
21    movl     $4, -8(%rbp)
22    movl     $4, -4(%rbp)
23    movl     -4(%rbp), %eax
24    popq     %rbp
25    ret
26 Leh_func_end1:
```

Above: original C code

Right: basic g++ assembly output

Below: optimized (O2) g++

assembly output



```
1 .section __TEXT,__text,regular
2 .globl _main
3 .align 4, 0x90
4 _main:
5 Leh_func_begin1:
6     pushq    %rbp
7 Ltmp0:
8     movq     %rsp, %rbp
9 Ltmp1:
10    movl     $4, %eax
11    popq     %rbp
12    ret
13 Leh_func_end1:
```

The difference with the -O2 flag is tremendous! The entire main function was a 17 instruction x86 segment that closely matched the source C code. There are definitely also some inefficiencies, for example at the very end, there are a lot of extraneous 4's being placed onto the stack and then eventually one is pushed into the return register and the function returns.

Conversely, on the optimized assembly, g++ did whatever it could to minimize the number of instructions required to run the output assembly, so it took the time to pre-process things like the return value. The main function is 5 instructions now, 4 of which are standard boilerplate like saving/restoring the base pointer and calling ret to return. The way I see it, the O2 flag gives g++ the directive to take extra time now to make it more efficient later, pre-compute anything that you already have enough information to compute and cut-out any unnecessary movement/storage of data.

Cout Function

The next type of optimization I explored was optimization of I/O using standard functions: namely cout. Although, I didn't expect to really be able to understand most of the generated assembly, as I imagine I/O to be complex on the lower levels, I did gain some valuable insights throughout reading through the assembled programs. I ended up reading a good deal into the differences between the standard x86-64 architecture and that of x86-64 Mac OS X computers.

The following are screen captures of the C++ source, the default assembly, and the optimized assembly (in that order):

```
Vim
1 #include <iostream>
2
3 using namespace std;
4
5 int main() {
6     cout << 127 << endl;
7 }
```

Original Source

Truncated assembly

```
Vim
1 .section __TEXT,__text,regular,pure_instructions
2 .globl _main
3 .align 4,0x90
4 _main:
5 Leh_func_begin1:
6     pushq %rbp
7     Ltmp0:
8     movq %rsp,%rbp
9     Ltmp1:
10    subq $16,%rsp
11    Ltmp2:
12    movq __ZSt4cout@GOTPCREL(%rip),%rax
13    leaq (%rax),%rax
14    movl $127,%ecx
15    movq %rax,%rdi
16    movl %ecx,%esi
17    callq __ZNSolsEi
18    movq __ZSt4endlC@GOTPCREL(%rip),%rcx
19    leaq (%rcx),%rcx
20    movq %rax,%rdi
21    movq %rcx,%rsi
22    callq __ZNSolsEPFRSoS_E
23    movl $0,-8(%rbp)
24    movl -8(%rbp),%eax
25    movl %eax,-4(%rbp)
26    movl -4(%rbp),%eax
27    addq $16,%rsp
28    popq %rbp
29    ret
30 Leh_func_end1:
31
32 .section __TEXT,__StaticInit,regular,pure_instructions
33 .align 4,0x90
34 __GLOBAL__I_main:
35 Leh_func_begin2:
36     pushq %rbp
37     Ltmp3:
38     movq %rsp,%rbp
39     Ltmp4:
40     movl $1,%eax
41     movl $65535,%ecx
42     movl %eax,%edi
43     movl %ecx,%esi
44     callq __Z41__static_initialization_and_destruction_0ii
45     popq %rbp
46     ret
47 Leh_func_end2:
48
49 .align 4,0x90
50 __Z41__static_initialization_and_destruction_0ii:
51 Leh_func_begin3:
52     pushq %rbp
53     Ltmp5:
54     movq %rsp,%rbp
55     Ltmp6:
```

Optimized Assembly

```
1  [section __TEXT,__text,regular,pure_instructions
2  .globl _main
3  .align 4, 0x90
4  _main:
5  Leh_func_begin1:
6  pushq %rbp
7  Ltmp0:
8  movq %rsp, %rbp
9  Ltmp1:
10 movq __ZSt4cout@GOTPCREL(%rip), %rdi
11 movl $127, %esi
12 callq __ZNSoIsEi
13 movq __ZSt4endlIcSt11char_traitsIcEERSt13basic_ostreamIT_T0_ES6_@GOTPCREL(%rip), %rsi
14 movq %rax, %rdi
15 callq __ZNSoIsEPFRSoS_E
16 xorl %eax, %eax
17 popq %rbp
18 ret
19 Leh_func_end1:
20
21 .section __TEXT,__StaticInit,regular,pure_instructions
22 .align 4, 0x90
23 __GLOBAL__I_main:
24 Leh_func_begin2:
25 pushq %rbp
26 Ltmp2:
27 movq %rsp, %rbp
28 Ltmp3:
29 leaq __ZStL8__ioinit(%rip), %rdi
30 callq __ZNSt8ios_base4InitC1Ev
31 leaq __tcf_0(%rip), %rdi
32 xorl %esi, %esi
33 movq __dso_handle@GOTPCREL(%rip), %rdx
34 popq %rbp
35 jmp __cxa_atexit # TAILCALL
36 Leh_func_end2:
37
38 .section __TEXT,__text,regular,pure_instructions
39 .align 4, 0x90
40 __tcf_0:
41 Leh_func_begin3:
42 pushq %rbp
43 Ltmp4:
44 movq %rsp, %rbp
45 Ltmp5:
46 leaq __ZStL8__ioinit(%rip), %rdi
47 popq %rbp
48 jmp __ZNSt8ios_base4InitD1Ev # TAILCALL
49 Leh_func_end3:
50
51 .zerofill __DATA,__bss,__ZStL8__ioinit,1,3
52 .section __DATA,__mod_init_func,mod_init_funcs
53 .align 3
54 .quad __GLOBAL__I_main
55 .section __TEXT,__eh_frame,coalesced,no_toc+strip_static_syms+live_support
56 EH_frame0:
57 Lsection_eh_frame:
```

The first thing I noticed is that the optimized assembly code passes much less to the `__ZNSoIsEi` function that is called. The original code passes in `__ZSt4cout@GOTPCREL($rip)` [via `$RAX`] to the function. This intrigued me, as it seemed like it was passing in a function to the function.

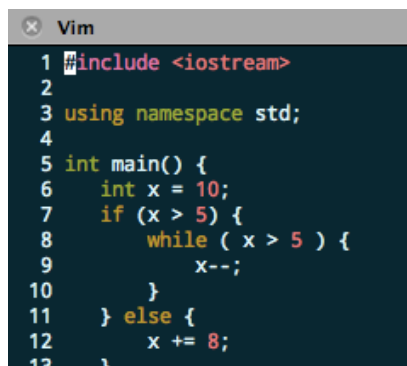
It turns out, the GOTPCREL(\$rip) is the Mac x86-64 environment's way of accessing local and small data. (Source: http://developer.apple.com/library/mac/documentation/DeveloperTools/Conceptual/MachOTopics/1-Articles/x86_64_code.html)

My interpretation of this is that the prefix `_ZSt4cout` is a label for one of the functions that exists in the `iostream` library (the `cout` function, perhaps a specific one to print a number vs. printing another data type) and is represented in the program space's "global offset table" (GOT) which uses "RIP-relative addressing" which means addressing relative to the instruction pointer.

The optimized code performs very similar things, but in fewer instructions. For example, the default code stores the `ZSt4cout` function pointer in `$rax` and then moves it from `$rax` to `$rdi`. This is probably due to convention. The optimized code on the other hand stores the information directly into `$rdi`. I don't understand why modern compilers don't always optimize little steps like these, as it doesn't even seem to be an optimization, but rather an improvement on the inefficient assembly produced by default. I'm sure there is a reason for this that involves things beyond the scope of my understanding.

Control Flow

The final optimization that I tried to study was that of control flow. I made a C++ program that included a very primitive demonstration of `if/else` and a `while` loop.



```
Vim
1 #include <iostream>
2
3 using namespace std;
4
5 int main() {
6     int x = 10;
7     if (x > 5) {
8         while ( x > 5 ) {
9             x--;
10        }
11    } else {
12        x += 8;
13    }
```

```

1  [section __TEXT,__text,
2  .globl _main
3  .align 4, 0x90
4  _main:
5  Leh_func_begin1:
6  pushq %rbp
7  Ltmp0:
8  movq %rsp, %rbp
9  Ltmp1:
10 movl $10, -12(%rbp)
11 movl -12(%rbp), %eax
12 cmpl $5, %eax
13 jle LBB1_5
14 jmp LBB1_3
15 LBB1_2:
16 movl -12(%rbp), %eax
17 subl $1, %eax
18 movl %eax, -12(%rbp)
19 LBB1_3:
20 movl -12(%rbp), %eax
21 cmpl $5, %eax
22 jg LBB1_2
23 jmp LBB1_6
24 LBB1_5:
25 movl -12(%rbp), %eax
26 addl $8, %eax
27 movl %eax, -12(%rbp)
28 LBB1_6:
29 movl $0, -8(%rbp)
30 movl -8(%rbp), %eax
31 movl %eax, -4(%rbp)
32 movl -4(%rbp), %eax
33 popq %rbp
34 ret
35 Leh_func_end1:
36

```

default

```

1  [section __TEXT,__text,regular,pure_instructions
2  .globl _main
3  .align 4, 0x90
4  _main:
5  Leh_func_begin1:
6  pushq %rbp
7  Ltmp0:
8  movq %rsp, %rbp
9  Ltmp1:
10 xorl %eax, %eax
11 popq %rbp
12 ret
13 Leh_func_end1:
14
15 [section __TEXT,__StaticInit,regular,pure_instructions
16 .align 4, 0x90
17 __GLOBAL__I_main:
18 Leh_func_begin2:
19 pushq %rbp
20 Ltmp2:
21 movq %rsp, %rbp
22 Ltmp3:
23 leaq __ZStL8_ioint(%rip), %rdi
24 callq __ZNSt8ios_base4InitC1Ev
25 leaq __tcf_0(%rip), %rdi
26 xorl %esi, %esi
27 movq __dso_handle@GOTPCREL(%rip), %rdx
28 popq %rbp
29 jmp __cxa_atexit # TAILCALL
30 Leh_func_end2:
31
32 [section __TEXT,__text,regular,pure_instructions
33 .align 4, 0x90
34 __tcf_0:
35 Leh_func_begin3:
36 pushq %rbp
37 Ltmp4:
38 movq %rsp, %rbp
39 Ltmp5:
40 leaq __ZStL8_ioint(%rip), %rdi
41 popq %rbp
42 jmp __ZNSt8ios_base4InitD1Ev # TAILCALL
43 Leh_func_end3:

```

optimized

The loop iteration in the non-optimized code is pretty similar to how I would code it myself. It creates a few labels for each of the control flow “states” and then uses the basic `cmp`, `jle`, `j`, `jge` instructions that I would use, were I to code this myself.

The optimized code on the other hand, is very different. It is much harder to draw similarities between the C++ code and the optimized assembly code. I tried to search for some common instructions like `cmp` or any jump, but to no avail, so I decided to dig around in the code and see what I could find through the Internet.

The first thing that jumped out to me was the line

```
jmp __cxa_atexit # TAILCALL
```

This seemed odd, because I hadn’t yet seen g++ generate comments in assembled code. It turned out that this was actually present in the default assembly code, as well as every other assembled program I had generated with

g++. It turned out to be g++'s way of doing constructors, so this must have been some sort of initialization of the main() object? This leads me to believe that this program simply does nothing, as the first (main) function simply xor's EAX with itself (zero's it) and then returns. The compiler rendered my function equivalent to "return 0;" and come to think of it, it technically was, as none of the variables or register values have any importance outside of the scope of the function itself.

Inheritance

In order to explore the assembly language that is created with classes, I created a simple "Numbers" class which holds three integer pointers, and then I created a subclass called "MoreNumbers" which holds an additional integer pointer. The code for the two classes are shown to the right and below:

```
5 class Numbers {
6 public:
7     Numbers( int *n1, int *n2, int *n3 );
8     ~Numbers();
9     int *x1;
10    int *x2;
11    int *x3;
12
13 };
14
15 Numbers::Numbers( int *n1, int *n2, int *n3 ) {
16     x1 = n1;
17     x2 = n2;
18     x3 = n3;
19 }
20
21 Numbers::~~Numbers() {
22     delete x1;
23     delete x2;
24     delete x3;
25 }
```

```
27 class MoreNumbers : public Numbers {
28 public:
29     MoreNumbers( int *n1, int *n2, int *n3, int *n4 );
30     ~MoreNumbers();
31     int *x4;
32 };
33
34 MoreNumbers::MoreNumbers( int *n1, int *n2, int *n3, int *n4 ) : Numbers ( n1, n2, n3 ) {
35     x4 = n4;
36 }
37
38 MoreNumbers::~~MoreNumbers() {
39     delete x4;
40 }
41 }
```

The main function of the program simply instantiates various integers, and creates one instance of Numbers and one of of MoreNumbers, and then deletes them in order to be able to properly examine the constructor and destructor of both a normal object and an object which inherits some of its data.

```

43 int main() {
44     int t1 = 0;
45     int t2 = 2;
46     int t3 = 3;
47     Numbers *n = new Numbers ( &t1, &t2, &t3 );
48
49     t1 = 5;
50     n->x1 = &t1;
51
52     t1 = 9;
53     t2 = 10;
54     t3 = 16;
55     int t4 = 15;
56
57     MoreNumbers *m = new MoreNumbers( &t1, &t2, &t3, &t4 );
58
59     delete n;
60     delete m;
61
62 }

```

To the left is a screen capture of the main function, showing its exact functionality, followed by the associated assembly code.

The first few lines in the snippet are

```

x Default
275 Ltmp29:
276 movl    $0, -44(%rbp)
277 movl    $2, -48(%rbp)
278 movl    $3, -52(%rbp)
279 movabsq $24, %rax
280 movq    %rax, %rdi
281 callq   __Znwmm
282 movq    %rax, -40(%rbp)
283 movq    -40(%rbp), %rax
284 leaq    -44(%rbp), %rcx
285 leaq    -48(%rbp), %rdx
286 leaq    -52(%rbp), %rsi
287 movq    %rax, %rdi
288 movq    %rsi, -88(%rbp)
289 movq    %rcx, %rsi
290 movq    %rdx, -96(%rbp)
291 movq    -88(%rbp), %rax
292 movq    %rcx, -104(%rbp)
293 movq    %rax, %rcx
294 callq   __ZN7NumbersC1EPiS0_
295 movq    -40(%rbp), %rax
296 movq    %rax, -64(%rbp)
297 movl    $5, -44(%rbp)
298 movq    -64(%rbp), %rax
299 movq    -104(%rbp), %rcx
300 movq    %rcx, (%rax)
301 movl    $9, -44(%rbp)
302 movl    $10, -48(%rbp)
303 movl    $16, -52(%rbp)
304 movl    $15, -68(%rbp)
305 movabsq $32, %rax
306 movq    %rax, %rdi
307 callq   __Znwmm
308 movq    %rax, -32(%rbp)
309 movq    -32(%rbp), %rax
310 leaq    -68(%rbp), %rcx
311 movq    %rax, %rdi
312 movq    -104(%rbp), %rsi
313 movq    -96(%rbp), %rdx
314 movq    -88(%rbp), %rax
315 movq    %rcx, -112(%rbp)
316 movq    %rax, %rcx
317 movq    -112(%rbp), %r8
318 callq   __ZN11MoreNumbersC1EPiS0_
319 movq    -32(%rbp), %rax
320 movq    %rax, -80(%rbp)
321 movq    -64(%rbp), %rax
322 movq    %rax, -24(%rbp)
323 movq    -24(%rbp), %rax
324 cmpq    $0, %rax
325 je      LBB11_2
326 movq    -24(%rbp), %rax
327 movq    %rax, %rdi
328 callq   __ZN7NumbersD1Ev
329 movq    -24(%rbp), %rax
330 movq    %rax, %rdi
331 callq   __ZdIPv

```

showing the three arguments for the

Numbers constructor (0, 2, 3) as written in the C++, but next there is an instruction that I have never seen before, `movabsq $24, %rax`.

From an online search, I found that this is meant to move an absolute quadword, and I think it may be passing in the size of the arguments ($24 = 8 * 3$, 8 for each pointer). The function that is called is called (Znwm) doesn't seem to be in any part of the assembly file, and it is called again later in the code to call the constructor for MoreNumbers, this is probably to do the inherited functionality from the base-class's constructor. On second thought, it appears to be a function that is called before the constructor, which I believe to be the function on line 294.

Next, it appears that the return value (the object itself?) is stored 40 bytes behind the base pointer. I know that 44 bytes behind the base pointer lies the t1 member variable, as line 297 seems to be setting it to five. This means that the first four

bytes (40-44 behind the base) are probably the implicit "this", and the following four byte chunks are the data members.


```

1  .section __TEXT,__text
2  .globl __ZN7NumbersC2EpiS
3  .align 1, 0x90
4  __ZN7NumbersC2EpiS0_S0_:
5  Leh_func_begin1:
6  pushq %rbp
7  Ltmp0:
8  movq %rsp, %rbp
9  Ltmp1:
10 movq %rdi, -8(%rbp)
11 movq %rsi, -16(%rbp)
12 movq %rdx, -24(%rbp)
13 movq %rcx, -32(%rbp)
14 movq -8(%rbp), %rax
15 movq -16(%rbp), %rcx
16 movq %rcx, (%rax)
17 movq -8(%rbp), %rax
18 movq -24(%rbp), %rcx
19 movq %rcx, 8(%rax)
20 movq -8(%rbp), %rax
21 movq -32(%rbp), %rcx
22 movq %rcx, 16(%rax)
23 popq %rbp
24 ret
25 Leh_func_end1:

```

Line 318 seems to be the call to the constructor of MoreNumbers, and in name alone, it seems to have a lot in common with the constructor of Numbers, except for the addition of “More” and one extra “S0_” suffix. This is probably a result of the inheritance. I was expecting the assembler to automatically call both constructors, perhaps calling the base-class’s constructor as a part of the call to the

constructor of the subclass, and I predicted correctly.

Above is the body of the Numbers constructor, and to the right is the body of the MoreNumbers constructor. As is apparent on line 93, the MoreNumbers constructor calls the Numbers constructor within itself. I cannot exactly tell how the data is manipulated after the constructor is called though.

```

70 .globl __ZN11MoreNumbersC2EpiS0_S0_
71 .align 1, 0x90
72 __ZN11MoreNumbersC2EpiS0_S0_:
73 Leh_func_begin4:
74 pushq %rbp
75 Ltmp6:
76 movq %rsp, %rbp
77 Ltmp7:
78 subq $48, %rsp
79 Ltmp8:
80 movq %rdi, -8(%rbp)
81 movq %rsi, -16(%rbp)
82 movq %rdx, -24(%rbp)
83 movq %rcx, -32(%rbp)
84 movq %r8, -40(%rbp)
85 movq -8(%rbp), %rax
86 movq -16(%rbp), %rcx
87 movq -24(%rbp), %rdx
88 movq -32(%rbp), %rsi
89 movq %rax, %rdi
90 movq %rsi, -48(%rbp)
91 movq %rcx, %rsi
92 movq -48(%rbp), %rcx
93 callq __ZN7NumbersC2EpiS0_S0_
94 movq -8(%rbp), %rax
95 movq -40(%rbp), %rcx
96 movq %rcx, 24(%rax)
97 addq $48, %rsp
98 popq %rbp
99 ret
100 Leh_func_end4:

```

```

323 je LBB11_2
326 movq -24(%rbp), %rax
327 movq %rax, %rdi
328 callq __ZN7NumbersD1Ev
329 movq -24(%rbp), %rax
330 movq %rax, %rdi
331 callq __ZdlPv
332 LBB11_2:
333 movq -80(%rbp), %rax
334 movq %rax, -16(%rbp)
335 movq -16(%rbp), %rax
336 cmpq $0, %rax
337 je LBB11_4
338 movq -16(%rbp), %rax
339 movq %rax, %rdi
340 callq __ZN11MoreNumbersD1Ev
341 movq -16(%rbp), %rax
342 movq %rax, %rdi
343 callq __ZdlPv
344 LBB11_4:
345 movl $0, -8(%rbp)
346 movl -8(%rbp), %eax
347 movl %eax, -4(%rbp)
348 movl -4(%rbp), %eax
349 addq $112, %rsp
350 popq %rbp
351 ret
352 Leh_func_end11:

```

I believe that per the x86 naming conventions of g++, the constructors are always the name of the class followed by a C, and the destructors with a D. That being said, the destructors are the next functions called in the main function (right, lines 328 and 340).

The prologue for the first destructor seems to be data passed in through the RDI register from behind the base pointer. This is probably the “this” pointer for the object that is allocated somewhere on the heap. The next function called at 331 “ZdlPv” is kind of confusing, but from purely the name it sounds like it **d**eletes a **p**ointer **v**ariable, and perhaps it is

responsible for deallocating the data members of the instance of the Numbers class. From searching through the Internet, I found references to this strange function on gnu.org and realized that it is a part of GCC and stands for "operator delete(void*)" (source: <http://lists.gnu.org/archive/html/bug-binutils/2010-01/msg00047.html>) meaning it deletes a void pointer (a point of any type). It seems to me like this is probably the function used to delete all pointers in every single destructor.

Predictably, the function is called again after the MoreNumbers destructor is called, leading me to believe that this Zdlpv is the function to actually free the memory itself, whereas the other functions that share the namesake of their respective classes exist simply to perform any other functionality that might need to exist for an objects destruction.